

# Evaluating Metagenome Assembly on a Complex Community

Sherine Awad<sup>1</sup>, Luiz Irber<sup>1</sup>, C. Titus Brown<sup>1,\*</sup>

**1 Population Health and Reproduction University of California, Davis, Davis, CA, USA**

**\* E-mail: ctbrown@ucdavis.edu**

## Abstract

### NEEDS ENHANCEMENTS

**Motivation:** With the emergence of de novo assembly, several work have been to done to assemble metagenomic data from de novo. Several assemblers exist that are based on different assembly techniques. However, we still lack a study that analyze different assemblers behavior on metagenomic data .

**Problem statement:** In this paper, we performed analytical study for metagnome assembly using different assemblers and different preprocessing treatments. The aim of the analysis is studying how well metagenome assembly works, and which assembly works best. In addition, the study analyzes the resource requirements of the assembly.

**Approach:** We used a mock community dataset for the analysis, and used its reference genome for benchmark evaluation. We quality filtered the reads, then we applied 2 other preprocessing steps: digital normalization and partitioning. We used 4 different assembler: Velvet, IDBA-UD, SPAdes, and MEGAHIT to assemble the reads using each treatment. We used QUAST to analyze assemblies accuracy.

**Results:** Results show that assembly works well. Velvet is the worst assembler in terms of accuracy and recourses utilizations. The results also showed that assembly counts to most of the reads.

**Conclusions:** Except for Velvet, assemblers works well. Further analysis is required to study which assembler is better used with each specific dataset. This step is left for our future work,

## Author Summary

WHAT SHOULD BE WRITTEN HERE

## Introduction

Metagenome is the sequencing of DNA in an environmental sample. While whole genome sequencing (WGS) usually targets one genome, metagenome targets several ones which introduces complexity to metagenome analysis due to genomic diversity and variable abundance within populations. Metagenomic assembly means the assembly of multiple genomes from mixed sequences of reads of multiple species in a microbial community. Most approaches for analyzing metagenomic data rely on mapping to reference genomes. However, not all microbial diversity of many environments are covered by reference databases. Hence, the need for de novo assembly of complex metagenomic data rises. Several assemblers exist that can be used for de novo assembly. In order to decide which assembly works best, we need to evaluate metagenome assembly generated by each assembler. In this paper, we provide, an evaluation for metegnome assembly generated by several assemblers and using different preprocessing treatments. We use a reference genome as a benchmark for the evaluation. The evaluation is based on assembly accuracy, and time and memory requirements. This evaluation shed light on doability of metagenome assembly and the minimum requirements needed for the assembly. In addition, knowing how each assembler works, helps deciding which assembler to use prior to assembly. However, the later point is left for our future work.

The comparative study in this paper is based on four different assemblers; Velvet [?], SPAdes [?], IDBA-UD [?], and MEGAHIT [?].

Velvet [?] is a group de Bruin graph-based sequence assembly methods for very short reads that can both remove errors. It also uses read pair information to resolve a large number of repeats. The

error correction algorithm merges the sequences that belongs together. Then the repeat solver algorithm separates parts that share overlaps.

SPAdes [?] is an assembler for both single-cell and standard (multicell) assembly. SPAdes generates single-cell assemblies and provides information about genomes of uncultivable bacteria that vastly exceeds what may be obtained via traditional metagenomics studies.

IDBA-UD [?] is a de Bruijn graph approach for assembling reads from single cell sequencing or metagenomic sequencing technologies with uneven sequencing depths. IDBA-UD uses multiple depth-relative thresholds to remove erroneous k-mers in both low-depth and high-depth regions. It also uses paired-end information to solve the branch problem of low-depth short repeat regions. It applies an error correction step to correct reads of high-depth regions that can be aligned to high confident contigs.

MEGAHIT [?] is a new approach that constructs a succinct de Bruijn graph using multiple k-mers, and uses a novel "mercy k-mer" approach that preserves low-abundance regions of reads. It also uses GPUs to accelerate the graph construction.

## Materials and Methods

### Datasets

We used a diverse mock community data set containing 64 known species, sequenced with Illumina HiSeq, yielding 109,629,496 paired-end sequences with an untrimmed length of 11.07 Gbp and an estimated insert size of  $\sim 380$  [?].

We received the reference genomes from the original authors (posted on j) and the original reads are available through the NCBI Sequence Read Archive at Accession SRX200676. Figure ?? shows the coverage profile of the reference genome, and the percentage of read with that coverage.

### Quality Filtering

We removed adapters with Trimmomatic v0.30 in paired-end mode with the Truseq adapters [?]. We next used the fastq\_quality\_filter from the FASTX-Toolkit v0.0.13.2 [?] to remove sequences using the parameters -Q33 -q 30 -p 50, which keeps all sequences with 50% or more bases with quality score greater than or equal to 30.

### Mapping

We aligned all quality-filtered reads to the reference metagenome with bwa aln (v0.7.7.r441) [?]. We aligned paired-end and orphaned reads separately using bwa aln samse. We then used samtools (v0.1.19) [?] to convert SAM files to BAM files for both paired-end and orphaned reads. To count the unaligned reads, we find the records with the "4" flag in the SAM files [?].

We found chimeric alignments with the bwa mem aligner using the default parameters (v0.7.7.r441). Chimeric alignments cut reads in two (or more). For each chimeric alignment, in the SAM file, there will be a primary alignment and at least one secondary alignment tagged SA. To count the chimeric alignments, we count the records with the "SA" flag in the SAM file [?].

To extract the reads that contribute to unaligned contigs, we mapped the quality filtered reads to the unaligned contigs using bwa aln (v0.7.7.r441) [?]. Then we used samtools [?] to retrieve the reads that are mapped to the unaligned contigs.

### Reference Coverage and Coverage Profile

To evaluate how much of the reference genome was contained in the read data, we used bwa aln to map reads to the reference genome. We then calculated how many reference bases are contained in at least one

mapped read (script `sam-calc-refcov-cmp.py`).

## Digital Normalization

We applied the `normalize-by-median.py` script from khmer v1.1 to execute abundance normalization (“digital normalization”, [?]) on the data, retaining paired reads and using a k-mer size of 20 (`-p -k 20`). We executed digital normalization with 4 hash tables, each 1 GB in size (`-N 4 -x 1e9`). After read normalization, we used the `filter-abund.py` script to trim high-abundance reads (estimated k-mer coverage  $\geq 20$ ) at low-abundance k-mers (k-mer abundance  $\leq 2$ ) to remove erroneous k-mers [?] [?].

## Partitioning

We next applied partitioning to the data [?,?]. We first eliminated high-abundance k-mers that could join multiple species bins using the `filter-below-abund.py` script from khmer v1.1 using an abundance cutoff of 50 or higher. We then ran `do-partition.py` with a k-mer size of 32 and 4 Bloom filters each of size 1 gigabit for partitioning (`-k 32 -N 4 -x 1e9`). After partitioning, partitions were extracted to groups using the `extract-partitions.py` script with a maximum group size of 100,000 (`-X 100000`).

## Metagenomes Assembly and evaluation

We assembled the reads using four different assemblers: Velvet [?], IDBA-UD [?], SPAdes [?], and MEGAHIT [?].

For Velvet v1.2.07 [?], we used k-mer values from 19 to 51 incremented by 2. We also used `-fastq.gz` for fastq format, `-shortPaired` for the pe files and `-short` for the se files. Also, we asked Velvet to automatically calculate expected coverage and coverage cutoff (`-exp_cov auto -cov_cutoff auto`). From among the many assemblies, one for each k-mer size, we then chose the assembly that had the most bases in contigs longer than 500 bp (script `calc-best-assembly.py`).

For IDBA-UD v1.1.1 [?], we used `-pre_correction` to perform pre-correction before assembly and `-r` for the pe files. For SPAdes v3.1.1 [?], we used `-sc -pe1-12` where `-sc` is required for MDA (single-cell) data and `-pe1-12` for file with interlaced reads for the first paired-end library.

For MEGAHIT [?], we used `-l 101 -m 3e9 -cpu-only` where `-l` is for maximum read length, `-m` is for max memory in byte to be used, and `-cpu-only` to use CPU not GPU.

We examined the assembly quality of each assembler and treatment using QUAST v2.3 [?] using `quast.py` and we use the default minimum contig length equal to 500.

## Results

### Cost of Assembly

We estimated time and memory requirements for each of them. We also estimated the running time and memory utilization for each assembler under both treatments and compared to assemblers time and memory requirements using quality filtered reads.

Digital normalization utilized 74.93 GB of memory and took around 3 hours and 53 minutes to run. Partitioning utilized 21.78 GB and around 2 hours and a half to run.

For Velvet assemblies, table ?? row 3, it took  $\sim 60$  hours using quality filtered reads, while it took only  $\sim 6$  hours using digital normalizations and  $\sim 4$  hours using partitioning. For IDBA-UD assemblies, table ?? row 6, it took  $\sim 33$  hours using quality filtered reads, while it took  $\sim 6$  hours using digital normalization and  $\sim 8$  hours using partitioning. SPAdes assemblies utilized  $\sim 67$  hours using quality filtered reads while it took  $\sim 15$  hours and  $\sim 7$  hours using digital normalization and partitioning respectively, table ?? row 9.

Finally, for MEGAHIT, it took  $\sim 2$  hours,  $\sim$  half an hour, and  $\sim$  hour and a half using quality-filtered reads, digital normalization, and partitioning respectively, table ?? row 12.

For Velvet assemblies, table ?? row 4, it used 98.40 GB of memory using quality filtered reads, while it used one 52.67 GB and 35.23 GB of memory when applying digital normalization and partitioning respectively. For IDBA-UD assemblies, table ?? row 7, it used used 123.84 GB of memory using quality filtered reads, while it used one 99.88 GB and 76.53 GB of memory when applying digital normalization and partitioning respectively. For SPAdes assemblies, table ?? row 10, it used 381.79 GB of memory using quality filtered reads, while it used one 121.52 GB and 94.70 GB of memory when applying digital normalization and partitioning respectively. For MEGAHIT, table ?? row 13 it utilizes 33.41 GB, 18.89 GB, 13.17 GB for quality-filtered reads, digital normalization, and partitioning respectively. See Table ?? for more details.

In terms of compute cost, Megahit is the fastest and utilizes less memory. In conclusion, quality filtering utilizes more resources digital normalization and the later utilizes more resources than partitioning.

## Assembly Comparisons

### Evaluation Against the Reference Genome

#### Assembly alignment

We aligned each assembly to the reference genome using nucmer cite here. The we used our script `analyze_assembly.py` to analyze the results. IDBA quality filtered assembly has 27,444 contigs, out of which 71.36 % is totally aligned to the reference genome. For SPAdes quality filtered assembly, it has 33,704 contigs and 72.22% is totally aligned to the reference. Finally, Megahit has 75,497 contigs and 74.97% is totally aligned to the reference.

#### Read Incorporation

To evaluate how much of the reads are captured by the assembly, we mapped the quality filtered reads to each assembly. Then we extracted the unaligned sequences.

#### More about unalignments

We extracted the reads that contributed to unaligned contigs of each assembler (see Methods:Mapping). We mapped those reads to the reference genome. We find that most of the reads that mapped to the unaligned contigs don't exist in the reference genome. Figure ?? shows a histogram for the percentage of reads contributed to unaligned contigs (orange columns) and the portion of those reads that mapped to the reference genome (green columns).

#### More about misassemblies

We extracted all the partially aligned contigs. Using IDBA and quality filtering, 11.99% of the contigs were partially aligned. 49.73% of the partially aligned contigs has less than 50 bases only unaligned while 50.27% has 50 bases or more unaligned. Using SPAdes and quality filtering, 10.82% of the contigs were partially aligned. 51.44% of the partially aligned contigs has less than 50 bases only unaligned while 48.56% has 50 bases or more unaligned. Using Megahit and quality filtering, 10.45% of the contigs were partially aligned. 58.49% of the partially aligned contigs has less than 50 bases only unaligned while 41.51% has 50 bases or more unaligned.

## More about uncovered regions

We used the same script to analyze the uncovered regions. Approximately 2.02% of the reference bases are uniquely covered by IDBA quality filtered assembly, ~9.75% and ~6.54% of the reference bases are uniquely covered by SPAdes and Megahit quality filtered assembly respectively using Besthit approach. Approximately ~10.60% of reference bases are uniquely uncovered by IDBA quality filtered assembly, ~7.62% and ~3.44% of the reference bases are uniquely uncovered by SPAdes and Megahit quality filtered assemblies respectively using Besthit approach. Meanwhile, there are 31,446,810 (~15.29%) of the reference bases commonly uncovered by IDBA, SPAdes, and Megahit quality filtered assemblies.

## Subsampling

Furthermore, we aligned the unaligned contigs of each assembly to the unaligned contigs of IDBA-UD quality filtered assembly. Using quality filtered reads, the reference fraction (the reference here is the unaligned contigs of IDBA-UD quality filtered assembly) equals to 80.613%, 91.922%, and 92.715% for Velvet, SPAdes, and MEGAHIT respectively. The unaligned length is 2,475,529, 2,174,574, and 916,247 for Velvet, SPAdes, and MEGAHIT respectively using quality filtered reads, representing 37.06%, 32.56% and 13.72% of the reference length. See Table ?? for more details. Velvet unalignments show the highest percentages of IDBA-UD QC unalignments. IDBA-UD diginorm unalignments shows the less percentage of IDBA-UD QC unalignments. This shows that not all reads are common among unaligned contigs.

## Discussion

### Assembly works well


Much of the reference is covered by the assembly. Most contigs are broadly accurate. (Q: How do we measure this?) Most genes within the reference are recovered. The assembly represents the majority of the reads (90%?) and the majority of the k-mers (XX).

## Acknowledgments


## Figure Legends

## Tables


## Supporting Information Legends




**Figure 1.** Total Length of assemblies in basepairs based on different min contigs length.




**Figure 2.** Genome Fraction of assemblies in basepairs based on different min contigs length.



**Figure 3.** *Genome Fraction of assemblies using different read coverage.*



**Figure 4.** *Reference genome coverage.*



**Figure 5.** Histogram for unaligned reads.



**Table 1.** Assembly Quality Metrics

<b>Treatment/Quality Metric</b>	<b>Quality Filtering</b>	<b>Digital Normalization</b>	<b>Partition</b>
<b>(1) Velvet</b>			
<b>Genome Fraction</b>	72.949	89.043	88.879
<b>Unaligned Length</b>	8,977,149	10,909,693	11,317,834
<b>Misassembled contigs length</b>	16,566,891	25,594,315	16,922,852
<b>N50</b>	38,028	18,944	8,504
<b>NG50</b>	22223	17212	7905
<b>(2) IDBA-UD</b>			
<b>Genome Fraction</b>	90.969	91.003	90.082
<b>Unaligned Length</b>	10,709,716	10,637,811	10,644,357
<b>Misassembled contigs length</b>	21,777,032	27,668,818	18,440,791
<b>N50</b>	49773	47828	26575
<b>NG50</b>	45748	44351	24326
<b>(3) SPAdes</b>			
<b>Genome Fraction</b>	90.424	90.173	89.272
<b>Unaligned Length</b>	10,597,529	10,621,398	10,500,235
<b>Misassembled contigs length</b>	28,238,787	23,103,154	14,338,099
<b>N50</b>	42773	35580	22319
<b>NG50</b>	38841	32598	19909
<b>(4) MEGAHIT</b>			
<b>Genome Fraction</b>	90.358	89.92	88.769
<b>Unaligned Length</b>	10,686,421	10,581,435	10,564,244
<b>Misassembled contigs length</b>	11,927,502	17,319,534	11,814,070
<b>N50</b>	35,136	27,302	17,492
<b>NG50</b>	32251	25248	15393

**Table 2.** Running Time and Memory Utilization

<b>Treatment/Quality Metric</b>	<b>Quality Filtering</b>	<b>Digital Normalization</b>	<b>Partition</b>
<b>(1) Velvet</b>			
<b>Running Time</b>	60:42:52	6:48:46	4:30:36
<b>Memory Utilization in GB</b>	98.40	52.67	35.23
<b>(2) IDBA-UD</b>			
<b>Running Time</b>	33:53:46	6:34:24	8:30:29
<b>Memory Utilization in GB</b>	123.84	99.88	89.25
<b>(3) SPADes</b>			
<b>Running Time</b>	67:02:16	15:53:10	7:54:26
<b>Memory Utilization in GB</b>	381.79	121.52	123.7
<b>(4) MEGAHIT</b>			
<b>Running Time</b>	1:52:55	0:30:23	1:23:28
<b>Memory Utilization in GB</b>	33.41	18.89	189.55

**Table 3.** mis-assemblies

Assembly	Quality Filtering	Digital Normalization	Partition
<b>(1) Velvet</b>			
mis-assemblies	917	5271	5202
Relocations	592	998	1036
Translocations	309	4262	4153
Inversions	16	11	13
Misassembled Contigs Length	16,566,891	25,594,315	16,922,852
Mismatches	104,740	174,446	178,348
Percentage of Mismatches	0.05%	0.08%	0.09%
Indels Length	50,190	181,453	346,988
Indels Percentage	0.02%	0.09%	0.17%
<b>(3) IDBA-UD</b>			
mis-assemblies	1223	1094	960
Relocations	613	668	578
Translocations	580	398	350
Inversions	30	28	32
Misassembled Contigs Length	21,777,032	27,668,818	18,440,791
Mismatches	162,733	231,432	230,840
Percentage of Mismatches	0.08%	0.11%	0.11%
Indels Length	30,433	43,358	42,523
Indels Percentage	0.01%	0.02%	0.02%
<b>(2) SPAdes</b>			
mis-assemblies	894	997	753
Relocations	608	613	496
Translocations	267	368	239
Inversions	19	16	18
Misassembled Contigs Length	28,238,787	23,103,154	14,338,099
Mismatches	184,630	244,849	235,396
Percentage of Mismatches	0.09%	0.12%	0.11%
Indels Length	27,328	32,783	21,516
Indels Percentage	0.01%	0.02%	0.01%
<b>(4) MEGAHIT</b>			
mis-assemblies	738	880	748
Relocations	448	593	513
Translocations	172	274	222
Inversions	118	13	13
Misassembled Contigs Length	11,927,502	17,319,534	11,814,070
Mismatches	152,964	207,349	203,515
Percentage of Mismatches	0.07%	0.10%	0.10%
Indels Length	15,298	18,195	16,517
Indels Percentage	0.01%	0.01%	0.01%

**Table 4.** Mapping quality-filtered reads to assemblies

Treatment	Quality Filtering	Digital Normalization	Partition
<b>(1) Velvet</b>			
No. of Unaligned Sequences	6,801,329	3,375,222	3,890,205
Percentage	6.40%	3.17%	3.66%
<b>(2) IDBA-UD</b>			
No. of Unaligned Sequences	1,490,609	1,738,371	2,297,377
Percentage	1.40%	1.63%	2.16%
<b>(3) SPAdes</b>			
No. of Unaligned Sequences	2,100,555	2,439,158	2,804,006
Percentage	1.98%	2.29%	2.64%
<b>(4) MEGAHIT</b>			
No. of Unaligned Sequences	1,559,300	2,082,881	2,747,427
Percentage	1.47%	1.96%	2.58%

**Table 5.** Mapping unaligned contigs to Idba quality-filtered unaligned contigs

Treatment/Quality Metric	Quality Filtering	Digital Normalization	Partition
<b>(1) Velvet</b>			
Unaligned IDBA Fraction	80.613 %	92.03%	92.982%
Unaligned Length	2,475,529	3192491	3,868,558
Percentage of unaligned	37.06%	47.8%	57.92%
<b>(2) IDBA-UD</b>			
Unaligned IDBA Fraction	-	91.53%	94.72%
Unaligned Length	-	498,299	1,320,036
Percentage of unaligned	-	7.46%	19.76%
<b>(3) SPAdes</b>			
Unaligned IDBA Fraction	91.92%	93.959%	94.826%
Unaligned Length	2,174,574	1,951,911	2,398,664
Percentage of unaligned	32.56%	29.22%	35.91%
<b>(4) MEGAHIT</b>			
Unaligned IDBA Fraction	92.715%	91.838%	92.219%
Unaligned Length	916,247	1,569,436	3,832,050
Percentage of unaligned	13.72%	23.5%	57.37%

**Table 6.** Supplementary Table: More Assembly Quality Metrics

<b>Treatment/Quality Metric</b>	<b>Quality Filtering</b>	<b>Digital Normalization</b>	<b>Partition</b>
<b>(1) Velvet</b>			
<b>N75</b>	13301	6084	3771
<b>NG75</b>	1186	4805	3214
<b>L50</b>	1013	2455	6037
<b>LG50</b>	1806	2740	6641
<b>L75</b>	2777	7026	14734
<b>LG75</b>	11087	8460	16867
<b>(2) IDBA-UD</b>			
<b>N75</b>	11693	12154	7834
<b>NG75</b>	9617	10221	6461
<b>L50</b>	828	896	1536
<b>LG50</b>	899	970	1712
<b>L75</b>	2986	3025	5062
<b>LG75</b>	3467	3484	6002
<b>(2) SPAdes</b>			
<b>N75</b>	11263	10554	6900
<b>NG75</b>	9005	8379	5401
<b>L50</b>	974	1192	1846
<b>LG50</b>	1078	1325	2108
<b>L75</b>	3276	3768	5840
<b>LG75</b>	3908	4495	7198
<b>(4) MEGAHIT</b>			
<b>N75</b>	8166	7230	5271
<b>NG75</b>	6601	5632	4030
<b>L50</b>	1199	1582	2490
<b>LG50</b>	1306	1757	2848
<b>L75</b>	4164	5063	7670
<b>LG75</b>	4907	6147	9581

**Table 7.** Comparision between N50 and NG50

<b>Treatment</b>	<b>Quality Filtering</b>	<b>Digital Normalization</b>	<b>Partition</b>
<b>(1) Velvet</b>			
<b>N50</b>	38,028	18,944	8,504
<b>NG50</b>	22,223	17,212	7,905
<b>(2) IDBA-UD</b>			
<b>N50</b>	49,773	47,828	26,575
<b>NG50</b>	45,748	44,351	24,326
<b>(3) SPAdes</b>			
<b>N50</b>	42,773	35,580	22,319
<b>NG50</b>	38,841	32,598	19,909
<b>(4) MEGAHIT</b>			
<b>N50</b>	35,136	27,302	17,492
<b>NG50</b>	32,251	25,248	15,393